

RQ4 Progress Report

Eamin C.

2026-09-06

Contents

Progress Report — RQ4	1
0. System diagram — the 7 sub-modules	1
1. Status of RQ4	3
2. What Component 3 needed to answer	3
3. The dataset at a glance	4
4. Three split strategies, three trade-offs	5
5. Why three strategies and not one	12
6. Context — Components 1 & 2	12
7. What is next	15
7.5 Live progress — 80-rollout pilot	15
8. How to reproduce	17


Progress Report — RQ4


For advisor review. Last updated: 2026-09-06 (80-rollout pilot live).

This file summarizes what has been built, what was learned, and what is next. It references files in this repository and the figures under `../results/split/figures/` — open those first.

0. System diagram — the 7 sub-modules

RQ4 is split into 7 sub-modules wired together as a one-way pipeline. Each stage consumes the artifacts of the previous stage and emits its own. Solid arrows = primary data flow; dashed arrows = control flow (prompts, configs).





Sub-module responsibilities

#	Module	Input	Output	Owner file
1	Question framing	literature	scope doc	README.md, README_HF_TOKEN_FIX.md
2	Issue collection + taxonomy	GitHub API + reading	200 labeled issues	utils/classify/, agent/, data/index.jsonl
3	Train/test split	data/index.json	16 indices × 3 sizes	utils/split/, utils/run_sweep.py
4	Per-repo skill training	train indices + train issues	agent/skills/<a>/<s>/<train_size>/	utils/train/train_skill.py
5	Verify pool	data/index.jsonl, is- + raw AgentSmith outputs	data/verify/_pool/, issue_index_*.jsonl	utils/verify/build_verify.py
6	Agent solve loop	verify index + skills	runs/<a>_<s>/.../patches/	utils/verify/solve.py
7	Score + plot + push	runs/.../patch.txt	results/rq4/*.csv, ut.png, ut.pdf	utils/verify/{score_patch, run_godpatch}

Why this shape

- **One-way** — each stage is reproducible from the artifacts of the previous one, so a bug at any stage only forces re-running stages 6 and 7 (the only stages that touch the LLM).
- **Repo-disjoint** between train (4) and verify (5) guarantees that module 6's f2p rate measures *true* skill transfer, not codebase memorisation.
- **Sanity check (7)** runs the gold patch through `run_f2p_verify` before scoring any agent — if the gold patch doesn't pass f2p on our infrastructure, the run is uninformative and is retried before scoring agents.

1. Status of RQ4

Component	Status	Artifact
1. Question framing + scope	done	README.md
2. Taxonomy	done (revised once)	docs/taxonomy.md (837 lines), taxonomy
2. Issue collection	done	data/issues/*.md, data/index.jsonl (200 issues)
2. Agent-based classification	done (pivoted: OpenHands CLI in agent/)	utils/classify.py, utils/run_classify.sh
3. Train/test split	done	utils/split/{run,dist_stats,visualize}, results/split/
4. Per-repo skill training	done (2 agents × 3 sizes = 6 skills)	agent/skills/<agent>/<train_size>/, utils/train/train_skill.py
5. Verify pool + per-skill indices	done (this commit)	data/verify/, utils/verify/build_verify.py
6. Agent solve loop on verify set	running (80-rollout pilot, 4-way parallel)	data/verify/runs/, utils/verify/solve.py
7. Score + plot + push	pending	—

This session covered Component 5 (verify pool). Components 1–4 are summarised briefly in §6 for context.

2. What Component 3 needed to answer

Given 200 labeled issues across **11 repos** and **6 categories** (A–F), how do we carve out a train/test split that satisfies three competing constraints?

1. **Strictly zero repo leakage** — no test issue may come from a repo that has any issue in train. Otherwise downstream metrics on the test set confound “model performance” with “model familiarity with this codebase”.
2. **Train size hits the requested N** — many experiment designs (e.g. few-shot prompting at N=40) need a specific count.
3. **Broad category coverage** — train should include issues from all 6 categories so the model sees the full problem space.

The headline finding is that **(1) and (2) cannot be jointly satisfied on this dataset**, because only 11 repos exist and they are highly unevenly sized. See §4.


3. The dataset at a glance



Source: data/index.jsonl □ results/split/distribution.md.

metric	value
total issues	200
distinct repos	11
distinct categories	6 (A, B, C, D, E, F)
largest repo	strands-agents/harness-sdk (79, 39.5%)
2nd largest	agentscope-ai/agentscope (50, 25.0%)
smallest	strands-agents/sdk-python (1, 0.5%)
repos with ≤ 5 issues	5 of 11
mean / median repo size	18.2 / 8

Category distribution: B(58) > C(38) > A(34) > E(33) > F(22) > D(15). D is the rare class — only 15 issues total.

Skewed, long-tailed. The top-1 repo contributes nearly 40 %; the bottom-5 together contribute ~5 %.





4. Three split strategies, three trade-offs

All three live in `utils/split/run.py` and are exposed via `--mode`. The sweep script (`utils/run_sweep.py`) was run for $\text{train_size} \in \{20, 40, 60, 80, 100, 120, 140, 160, 180\} \times 3 \text{ seeds} \times 3 \text{ modes} = 81 \text{ configurations}$, written to `results/split/sweep.{jsonl,md}`.

A. default — bounded per-repo cap (best-effort)

1. Compute per-category seed quotas via largest-remainder rounding.
2. Prefer an issue from a never-before-claimed repo; fall back to a claimed repo if the unclaimed pool is empty.
3. After seeding, top up each claimed repo to a per-repo cap $m_R = \lfloor \text{train_size} / |\text{claimed_repos}| \rfloor$ (rounded).

Effect: covers every repo and every category. Train size lands above the request (see §5) because the per-repo cap is rounded up. Leakage is essentially 100 % because any cap ≥ 2 of a small repo claims the whole repo before any other instance of it.

B. repo_disjoint — strict whole-repo isolation

1. Sort repos ascending by issue count.
2. Pick the smallest number N of repos such that $\sum \text{issue_counts}[r] \geq \text{train_size}$.
3. Claim **every** issue from those N repos; everything else $\square$ test.

Effect: by construction $\text{test_repo_leakage_pct} == 0$. Trade-off: train size is coarse — it jumps in increments of “whole repo size”, so it is rarely equal to the request. On this dataset it shows two plateaus: a step from $k=9 \rightarrow 10$ jumps $50 \square 121$ (one large repo enters the train set), and $k=10 \rightarrow 11$ jumps $121 \square 200$ (the whole corpus).

Note (bug fix): an earlier version of step 2 used a lexicographic sort plus `rng.shuffle`, which made the chosen k depend on the seed in a way that broke monotonicity (e.g. $\text{req}=40$ gave a *larger* train set than $\text{req}=80$ for some seeds). The ascending-by-size sort makes step 2 a pure function of `train_size`, so every seed now returns the same train/test sizes and the actual train is monotone non-decreasing in the request.

C. greedy_issue — issue-level greedy

1. Pass 1: snake through the 6×11 (category $\times$ repo) cross-tab; take at most one issue per cell.
2. Pass 2: top up to `train_size` from the smallest-claimed-repo pool.

Effect: hits `train_size` almost exactly (see §5). Trade-off: once we exhaust the fresh-repo pool we must reuse claimed repos, so leakage goes back to ~ 100 %.

Theoretical analysis: a probability / combinatorics framing

We formalise the split problem as follows.

Setup. Let R be the set of repos, with sizes $r_i = |\text{repo}_i|$ and $\sum_i r_i = N$. Let C be the set of categories. Define the counts matrix n_{rc} = number of issues of category c in repo r . Let $k \in [1, N)$ be the requested training set size. Let $R_{\text{train}} \subseteq R$ be the set of repos that contribute at least one issue to the training set. Define the *repo leakage fraction* as

$$L = \frac{|\{i \in \text{test} : \text{repo}(i) \in R_{\text{train}}\}|}{N_{\text{test}}} \quad (L = 0 \text{ strict isolation, } L = 1 \text{ full contamination}).$$

Strategy A — Bounded Per-Repo Cap A claims every repo (because the per-repo cap is ≥ 1 for all repos when $k \leq N$). Therefore $R_{\text{train}} = R$ always, and

$$L_A^{\min} = L_A^{\max} = 1, \quad L_A^{\text{actual}} \approx 1.0 \quad (100\% \text{ in all sweeps, } k = 20, \dots, 180).$$

The expected train size is exactly k ; the variance across seeds arises only from the order in which categories and repos are visited. A is **deterministically non-leakage-free** — no amount of randomisation changes L_A .

Strategy B — Repo-Disjoint (Strict Whole-Repo Isolation) Sort repos by size ascending: $r_{(1)} \leq r_{(2)} \leq \dots \leq r_{(R)}$. Define the prefix sums

$$S_k = \sum_{i=1}^k r_{(i)}, \quad S_0 = 0.$$

B claims exactly the smallest k^* repos where

$$k^* = \min\{k : S_k \geq k\}. \quad (1)$$

The actual training set size is S_{k^*} (the sum of those repos). Since repos are taken whole, $k \leq S_{k^*} < k + r_{(k^*+1)}$, so the approximation error is bounded by the size of the *next* repo.

Best case (most efficient packing). The repo sizes are tiny relative to k , so $k^* \approx k/\bar{r}$ and the excess $S_{k^*} - k$ is small. The leakage is strictly

$$L_B = 0 \quad \text{by construction.} \quad (2)$$

Worst case (least efficient packing). A single dominant repo covers almost the whole corpus, e.g. $r_{(R)} \approx N$. Then any $k < r_{(R)}$ forces $k^* = R - 1$ and $S_{k^*} = N - r_{(R)} \approx 0$, so the training set may be much smaller than k . The bound on under-fill is

$$k - S_{k^*-1} < r_{(k^*)}. \quad (3)$$

i.e. the worst-case gap is exactly the size of the first repo that *does* reach the threshold. For our dataset the 10th repo (79 issues) causes the $\text{req} = 80 \rightarrow \text{train} = 121$ plateau; the gap is $121 - 80 = 41$, which equals $r_{(11)} = 79$ capped by the last partial sum.

Theoretical summary for B:

quantity	formula	dataset value
leakage	0 (always)	0.0%
best-case train	k	never achieved
worst-case train	S_{k^*} in (1)	see plateau table
error bound	$< r_{(k^*+1)}$	max gap = 79
monotonicity	S_k monotone $\uparrow$	confirmed

Strategy C — Greedy Issue-Level C makes two passes over the $R \times C$ category–repo matrix.

Pass 1 (fresh repos): each of the $R \times C$ cells contributes at most one issue to train, drawn from an unclaimed repo. Let

$$F = \sum_{r \in R_{\text{train}}} \sum_{c \in C} \min(1, n_{rc}) \quad (\text{issues from fresh cells}).$$

Pass 2 (repo reuse): if $F < k$, C reuses already-claimed repos until the size reaches k . All reused issues introduce leakage.

Best case: $k \leq F$ — the request is satisfied entirely in Pass 1, no repo is reused, and $L_C = 0$. This requires $k \leq R \times C = 11 \times 6 = 66$ in general, or $k \leq 66$ on our dataset (confirmed by sweep: $k = 20, 40, 60$ all show leakage $< 100\%$, with $k = 60$ at the boundary).

Worst case: $k > F$ — Pass 1 is exhausted and Pass 2 must reuse every claimed repo. Every test issue shares a repo with train, so

$$L_C^{\max} = 1 \quad (\text{full contamination}). \quad (4)$$

Theorem (strategic trade-off). For any dataset with R repos, C categories, and N total issues:

strategy	leakage	train size	other property
A	$L_A \in \{1\}$	exact hit on k	no isolation guarantee
B	$L_B = 0$	bounded error $< r_{(k^*+1)}$	monotone, but coarse
C	$L_C \in [0, 1]$	exact hit on k	isolation until $k > F$

No strategy can simultaneously achieve $L = 0$ and $\text{train} = k$ unless $|R|$ is large enough (in our case $R = 11$) or the request k is small enough that the smallest repos suffice. This fundamental tension is the reason we recommend **B** for RQ4’s final reported experiment: it provides a *guarantee*, not just an empirical observation.

Concrete numbers on this dataset $N = 200$, $R = 11$, $C = 6$. Sorted repo sizes

$$r_{(1..11)} = [1, 2, 2, 3, 5, 8, 11, 15, 24, 50, 79],$$

with prefix sums

$$S_k = [1, 3, 5, 8, 13, 21, 32, 47, 71, 121, 200].$$

Strategy A (per-repo cap $m_R = \lfloor k/| \text{claimed} | \rfloor \approx k/11$):

k	per-repo cap	actual train	leakage
20	$\lfloor 20/11 \rfloor = 1$	28.7	99.2 %
60	$\lfloor 60/11 \rfloor = 5$	76.7	100 %
100	$\lfloor 100/11 \rfloor = 9$	120.3	100 %
180	$\lfloor 180/11 \rfloor = 16$	184.3	100 %

In every row $|R_{\text{train}}| = 11 = R$, confirming $L_A = 1$ from the preceding theorem. Excess over request is the price of rounding the per-repo cap up so each small repo contributes at least one issue.

Strategy B (closed form using prefix sums above):

k	$k^* = \arg \min S_k \geq k$	actual train = S_{k^*}	error $S_{k^*} - k$
20	6	21	+1
40	8	47	+7
60	9	71	+11
80	10	121	+41 (plateau)
100	10	121	+21
120	10	121	+1
140	11	200	+60
160	11	200	+40
180	11	200	+20

The +41 plateau at req = 80 is exactly the jump from $S_9 = 71$ to $S_{10} = 121$ (adding the 50-issue repo) — i.e. inequality (3) holds with the largest gap $41 < r_{(11)} = 79$. For req ≥ 140 , k^* saturates at $R = 11$ and train collapses to the whole corpus, leaving test empty.

Strategy C (fresh-cell bound $F \leq R \times C = 66$):



k	F achievable?	leakage
20	yes (well below 66)	98.7 %
60	yes (boundary)	100 %
80	no (Pass 2 forced)	100 %
180	no	100 %


The 98.7 % at $k = 20$ (not 100 %) is the only case where Pass 1 still dominates: the snake through the 66 (repo, category) cells picks 21 distinct cells, and Pass 2 reuses at most 1 issue — most test issues are from repos that were never claimed, hence $L < 1$. From $k \geq 60$ onwards Pass 1 is exhausted and $L_C = 1$.

Numerical comparison (averaged over 3 seeds)


Source: results/split/sweep.md, side-by-side table.


requested	A train	A leakage	B train	B leakage	C train	C leakage
20	28.7	99.2 %	21.0	0.0 %	20.3	98.7 %
40	53.3	100.0 %	47.0	0.0 %	43.0	100.0 %
60	76.7	100.0 %	71.0	0.0 %	62.7	100.0 %
80	97.7	100.0 %	121.0	0.0 %	83.7	100.0 %
100	120.3	100.0 %	121.0	0.0 %	106.0	100.0 %
120	138.3	100.0 %	121.0	0.0 %	128.0	100.0 %
140	153.0	100.0 %	200.0	0.0 %	147.7	100.0 %
160	167.7	100.0 %	200.0	0.0 %	168.0	100.0 %
180	184.3	100.0 %	200.0	0.0 %	188.0	100.0 %





Strategy comparison — averages over 3 repeats





5. Why three strategies and not one

The structural bottleneck is **(a) only 11 repos** and **(b) highly uneven sizes**. Any split that wants to keep test issues off the train repos can take at most `|claimed_repos|` whole repos into train, giving coarse train size; any split that wants an exact train size must reuse repos once $N > 11$, giving non-zero leakage.

Recommendation for RQ4:

- For the final, quoted experiment, use **B (repo_disjoint)**. Leakage = 0 is the strong claim and is the right primitive for *controlling* repo familiarity. Do not try to hit an exact N.
- For the few-shot prompt-size ablation ($N=20, 40, 80, 160 \dots$), declare the actual train/test sizes as part of each result cell — they are coarse but reproducible.
- For exploratory runs where exact N matters, **C (greedy_issue)** is fine but treat leakage $\approx 100\%$ as a known caveat.

If a future run demands both exact N and 0 leakage, the only way is to either (i) collect more repos / issues, or (ii) accept a leakage budget between 0 and 100 % by limiting train to $k \leq 2$ issues per repo and reporting the residual leakage as $1 - |\text{unique train repos}| / |\text{total repos}|$.

6. Context — Components 1 & 2

6.1 Component 1 — scope and question

README.md. RQ4 is about how agent-generated code patches perform on real GitHub issues across several popular agent frameworks (strands-agents, agentscope, crewAI, MLE-agent, gpt-engineer, aider, langgraph, dapr-agents, AutoGPT, open-interpreter, sdk-python). 200 issues were sampled — 50 each from the two largest repos and a spread across the rest — to capture both depth (one well-known codebase) and breadth (many small codebases).

6.2 Component 2 — taxonomy

docs/taxonomy.md (837 lines). 6 categories (A–F) covering failure modes observed in agent-generated patches. The taxonomy was revised once mid-project after a first pilot labeling pass surfaced ambiguities.

Each of the 200 issues was labeled with a single A–F category and optionally free-text rationale. The classifier was an OpenHands CLI agent (`utils/classify.py`, `utils/run_classify.sh`), not a hosted LLM. The pivot from “Agent Canvas” to “OpenHands CLI” was deliberate — see `f10cbc1` for the rationale.

6.3 Component 4 — per-repo skill training

`utils/train/train_skill.py` (bash `utils/train/run_all.sh` runs all 6 skills end-to-end). For each (agent, train_size) pair, we let the agent look at *every* train issue from a claimed repo and write a per-repo summary `agent/skills/<agent>/<train_size>/repos/<owner>__<name>.md`. All repo-level summaries are then concatenated with a template (`prompts.SKILL_TEMPLATE`) into the agent’s system prompt via the `SKILL.md` file in the same directory. The same template is also used to write a generic fallback `generic_fix.md` for repos the agent *did not* see at train time (e.g. test repos).

Six skills were trained:

```
agent/skills/  
├─ mini-swe-agent/{40,60,80}/  
└─ openhands/      {40,60,80}/
```

Each skill ships with a `manifest.json` recording provenance (which issues were used as train, which LLM was called, what cost was incurred). A `single_issue / single_repo` smoke-test mode lets us verify the pipeline end-to-end on one issue before paying for the full sweep.

6.4 Component 5 — verify pool (this commit)

`utils/verify/build_verify.py` writes two artefacts:

1. **data/verify/_pool/<repo-owner>__<repo-name>__<issue-id>/** — a *patch-stripped* copy of every raw issue directory (`data/raw/results/all_combined_f2p/...`). **200 dirs** (one per (repo, id) row in `data/index.jsonl`), **~4.8 MB total** (under `--strict`, the default), exactly **3 files per directory**:
 - `env.dockerfile` — kept verbatim.
 - `agentsmith_fail2pass_<NNN>.py` — the f2p test, kept verbatim (filename varies by issue).
 - `issue.json` — rewritten with the gold patch removed: `linked_prs[].patch`, `linked_prs[].base_sha`, `linked_prs[].head_sha` are stripped. Public PR metadata (number, state, title, url, merged, base_branch) is preserved. The pool writer additionally re-stamps repo and id so each dir is self-describing. The **repo-prefixed dir naming** (`<owner>__<name>__<id>`) is deliberate: GitHub issue numbers are repo-scoped, and the same number in two different repos refers to two different bugs. The previous id-only scheme silently overwrote content for the 8 colliding ids (e.g. `issue-563` in `AntonOsika/gpt-engineer` vs `dapr/dapr-agents`), so 16 index rows pointed at the wrong pool content. The audit catches any regression under the `pool_dir_uniqueness` check.
 - **generated_patch.diff** is never copied (115 raw dirs had it). `summary.json` / `agent-smith_stat.json` / `run.log` / `f2p.txt` / `dockerbuild.txt` are unconditionally dropped under `--strict`; they are only kept under the legacy `--no-strict` mode for forensic debugging of the pipeline itself. The strict 3-file invariant is the *only* contract the release audit checks against; downstream `solve.py` (Component 6) is written against it.
2. **Six per-skill indices**, `data/verify/issue_index_<agent>_<train_size>.jsonl`, one row per issue (200 rows each), `split ∈ {0, 1}`:
 - 0 = train issue, the agent must NOT be shown this when solving (skill contamination);
 - 1 = test issue, the agent is shown this and asked to produce a patch.

All 6 indices share the same `seed=42`, `mode=repo_disjoint` so the issue-level train/test partition is identical across skills; only the *size* of train (and hence test) varies. The split assignment is keyed by (repo, id) to avoid a latent dedup bug in `utils/split/run.py` where 8 ids that collide across repos (`issue-563`, `issue-974`, ...) used to silently vanish from test. The verify indices themselves

are keyed the same way, and the `verify_dir` field on each row points at the repo-prefixed pool dir (e.g. `data/verify/_pool/strands-agents__harness-sdk__issue-1702/`).

skill	train	test	repo leak
mini-swe-agent_40 / openhands_40	47	153	0 %
mini-swe-agent_60 / openhands_60	71	129	0 %
mini-swe-agent_80 / openhands_80	121	79	0 %

Monotonicity check: any issue in train at `train_size=40` is also in train at `train_size=60` and `80` (0 violations across the 200 issues). Same-size / cross-agent agreement: 0 mismatches.

A `verify_manifest.json` records the per-split counts and a leak audit (`audit.json`) reports 200/200 issues clean plus a `_GLOBAL_.pool_dir_uniqueness` check — every entry passes absent:generated_patch.diff, issue_json_no_patch_keys, **and the strict_3_file_invariant** (exactly `env.dockerfile` + `issue.json` + 1 `agentsmith_fail2pass_*.py`, nothing else). The uniqueness check verifies that every pool dir's `issue.json.{repo, id}` matches the `(repo, id)` row in `data/index.jsonl` — i.e. no two rows share a pool dir, and no row points at content belonging to a different `(repo, id)`. Run again any time with `python utils/verify/build_verify.py --audit` (idempotent; `--force-pool` rebuilds the pool from scratch, accepting the default `--strict 3-file` mode; pass `--no-strict` to switch back to the legacy 7-file mode for debugging the pipeline itself).

6.5 Component 6 — solver (scaffolding landed)

`utils/verify/solve.py` (521 lines, freshly added) drives the agent over the verify pool. Three sub-commands:

Sub	Status	Notes
<code>solve.py dry-run</code>	Landable now.	No docker / no LLM. Resolves SKILL.md (or <code>fallback_generic_fix.md</code>), builds the prompt, prints every input.
<code>solve.py run</code>	Helpers complete, branch stub.	<code>build_image_for_row</code> , <code>run_mini_in_docker</code> , <code>build_prompt</code> are fully wired. The CLI branch invokes <code>build_image_for_row</code> for the pilot row but stops short of <code>run_mini_in_docker</code> on the canonical pilot — see §7 for blockers.
<code>solve.py score</code>	TODO.	Apply trajectory's <code>patch.txt</code> to a fresh container, run the f2p test, record pass/fail.

Each `issue_index` row now also carries `base_sha` (read from the raw issue's first linked PR). 200/200 rows have it. Solver uses it to `git checkout` the upstream repo clone before `docker build`.

Dry-run pilot on **issue-1077** (the canonical strands-agents/ harness-sdk bug With bedrock guardrails, tool output is redacted breaking the conversation):

```

skill_key = mini-swe-agent_40
with-skill    prompt = 9 805 chars  (skill: 2 959 chars)
no-skill      prompt = 6 827 chars  (fallback: 1 815 chars)
verify_dir    = data/verify/_pool/strands-agents__harness-sdk__issue-1077/
pool_files    = [env.dockerfile, issue.json, agentsmith_fail2pass_1077.py]
base_sha      = 95906faf85095af9438a9bad072d437fd49b70e6

```

This confirms: pool is correct, prompt assembly is correct, and the with-skill variant contains the SKILL.md body while the no-skill variant contains only the generic fallback. See docs/handoff.md §4.1 for the full pilot-flow description.

The next step (§7) is to drive the agent on every split=1 issue in each of the 6 indices, with and without the corresponding SKILL.md injected, and record pass/fail.

7. What is next

1. **Drive the solve loop on every test issue in every skill index** (Component 6). For each of the 6 per-skill indices, iterate split=1 rows, point the agent at data/verify/_pool/<owner>__<name>__<id>/, inject the matching agent/skills/<agent>/<train_size>/SKILL.md into the system prompt, and record pass/fail against the agentsmith_fail2pass_<NNN>.py test. Also run a *no-skill* baseline (same issue, same agent, SKILL.md removed) so the per-skill lift can be reported. Expected cost: $\sim 6 \times 153$ issues per skill (lowest-size split) $\square$ ~ 918 runs with skill + 918 without = $\sim 1.8k$ runs.
2. **Reduce to the LLM metrics the paper needs:** pass rate per-skill, per-category (A–F), per-repo, lift from skill vs baseline, and statistical confidence (95 % CI via Wilson interval). Render all of this into a 1-page summary figure alongside the existing results/split/figures/.
3. **Open question: how to score partial credit?** Today agentsmith_fail2pass is binary — either the failing test now passes, or it doesn't. For agent patches that “look right” but don't flip the test, we currently mark zero. Worth considering a second signal (e.g. embedding-similarity to the gold patch) so we don't throw away useful signal.
4. **Add box plots** to utils/split/visualize.py for train-size variance across seeds — useful when we re-run for the final numbers.

7.5 Live progress — 80-rollout pilot

This section is updated live as the 80-rollout pilot (started 2026-09-06 ~22:34 UTC) progresses. Each rollout is one (issue, agent, skill_mode) triple. There are 20 issues $\times$ 2 agents $\times$ 2 skill_modes = 80 rollouts. The driver is utils/verify/run_pilot_20.py with 4-way concurrency.

Pilot issue selection: 20 test issues (split=1) sampled repo-diverse from data/verify/issue_index_mini-swe-agent_40.jsonl. Same 20 used for both agents. See data/verify/pilot_20_issues.jsonl.

Patch format quirks: - mini-swe-agent returns its submission text (a --- X.bak / +++ X unified diff). solve.py:extract_patch rewrites it to git format (inserts diff --git, drops .bak, strips /app/). - openhands returns nothing structured; solve.py:run_one_agent now cds into the testbed clone before invoking the CLI so the agent's git diff > patch.txt captures the real repo diff. After the agent finishes, the host's patch.txt is copied back to spec.out_dir.

Final result (updated 2026-09-07 01:30 UTC)

The batch driver completed but only **13 of 80** rollouts produced a patch.txt. The remaining 67 failures are all in the same upstream repo (strands-agents/harness-sdk) — its env.dockerfile in the verify pool uses hatch-vcs which fails on the upstream's git tag layout (Can't parse version from tag

'python/v1.14.0'). The build is broken unconditionally; not a per-agent issue. This is logged in each issue's `image_meta.json` with `ok=false`.

Restricting to the 13 buildable rollouts (4 issues × 2 agents × ~2 skill modes):

Combo	n	f2p	p2p	p2f	f2f	error	pass@1
mini + with-skill	3	0	1	0	1	1	0%
mini + without-skill	3	0	0	0	1	2	0%
openhands + with-skill	4	0	0	0	3	1	0%
openhands + without-skill	3	0	0	0	2	1	0%

Figures: - `results/rq4/figures/pilot20_outcomes.png` — all 80 rollouts (mostly `no_patch`) - `results/rq4/figures/pilot20_passrate.png` — buildable only - `results/rq4/figures/pilot20_breakdown.txt` — text breakdown

Honest read

- `pass@1 = 0` across all four combos. The 4 buildable issues (1 from `crewAI`, 3 from `agentscope`) are *not* trivial, and both agents over-engineered patches that either wouldn't apply (mini-1439 patch was in mini's submission format, my git-format rewrite tripped on a malformed `----` line), introduced syntax errors (oh-1439 patch had `i18n: I18N`, floats outside a function), or made no-op edits (mini-102 p2p — test passed both before and after, indicating the agent didn't reach the failing path).
- The `f2p` column was never populated. This is consistent with both agents being weak at structured-edit tasks on small repos; neither reliably converged on the gold-patch's minimal-diff shape.
- `strands-agents` is a separate problem: even if the agents had been perfect, those 67 rollouts would still report `no_patch` because the `verify-pool` docker image never built. Fixing that needs either a patched dockerfile or a different package-versioning scheme. Recorded as **known-infrastructure-issue** for the next iteration.

The full summary JSON (`results/rq4/pilot_20_summary.json`) is written when the batch driver finishes. Per-rollout scores are in `results/rq4/pilot_20_scores.csv`.

Code changes this session

- `utils/verify/solve.py`
 - added per-repo `fcntl.flock` around `build_one_image` to stop concurrent workers from racing on `.git/index.lock`.
 - added cleanup of stale `.git/index.lock` before each build.
 - `run_one_agent` now `cds` into the testbed clone for openhands so the agent's `git diff > patch.txt` captures the real repo diff and is then copied back to `spec.out_dir`.
 - added “Tools preference (CRITICAL)” to the prompt template, telling openhands to use the `bash` tool (avoiding the runtime's `file_editor` schema bug) and to write `patch.txt` from the repo root, not `/testbed`.
 - `extract_patch` for openhands now prefers the host-side `patch.txt` already copied by `run_one_agent` before falling back to `stdout` scraping.
- `utils/verify/run_pilot_20.py` (new) — 4-way concurrent driver over a JSONL issue list.
- `utils/verify/score_all.py` (new) — scores every rollout in the JSONL via `score.py` `score`, aggregates into `pilot_20_scores.csv`.
- `utils/verify/plot_pilot20.py` (new) — generates the two PNGs and the breakdown text from the CSV.
- `docs/progress.md` — added §0 system diagram (7 sub-modules pipeline) and §7.5 live-progress section.

8. How to reproduce

1. Generate distribution stats over the 200 issues

```
python utils/split/dist_stats.py \  
  --index data/index.jsonl --out results/split
```

2. Sweep train_size × seed × mode

```
python utils/run_sweep.py \  
  --index data/index.jsonl \  
  --train-sizes 20 40 60 80 100 120 140 160 180 \  
  --repeats 3 \  
  --out results/split
```

3. Render figures

```
python utils/split/visualize.py
```

4. Materialise one concrete split, e.g. mode B at train_size=60

```
python utils/split/run.py \  
  --index data/index.jsonl \  
  --train-size 60 \  
  --mode repo_disjoint \  
  --out results/split/final
```

5. Train 6 skills (one LLM call per repo per skill)

```
bash utils/train/run_all.sh
```

6. Build the verify pool + 6 per-skill indices

```
python utils/verify/build_verify.py          # idempotent (--strict default)  
python utils/verify/build_verify.py --audit  # + leak-vector scan + 3-file invariant  
# forensic: rebuild the legacy 7-file pool (only for debugging the pipeline)  
python utils/verify/build_verify.py --force-pool --no-strict --audit
```

7. (next) Run the solve loop on every test issue in every skill

```
python utils/verify/solve.py \  
  --agent openhands --train-size 40 \  
  --index data/verify/issue_indexopenhands_40.jsonl \  
  --split test \  
  --out results/verify/openhands_40.jsonl
```

This produces train.jsonl, test.jsonl, summary.json under results/split/final/, the full agent/skills/tree, and data/verify/{_pool, issue_index_*.jsonl, verify_manifest.json, audit.json}.